



# COG

CLOUD OPTIMIZED  
**GEOTIFF**

*Enabling Efficient Cloud Workflows*



## DOOM Emacs

```
87654321 0011 2233 4455 6677 8899 aabb ccdd eeff 0123456789abcdef
00000000: 4949 2a00 0800 0000 0f00 0001 0300 0100 II*.....
00000010: 0000 9702 0000 0101 0300 0100 0000 1b05 .....
00000020: 0000 0201 0300 0100 0000 4000 0000 0301 .....@.....
00000030: 0300 0100 0000 0100 0000 0601 0300 0100 .....
00000040: 0000 0100 0000 1101 0400 1b05 0000 2e15 .....
00000050: 0000 1501 0300 0100 0000 0100 0000 1601 .....
00000060: 0300 0100 0000 0100 0000 1701 0400 1b05 .....
00000070: 0000 c200 0000 1c01 0300 0100 0000 0100 .....
00000080: 0000 5301 0300 0100 0000 0300 0000 d885 ..S.....
00000090: 0c00 1000 0000 9a29 0000 af87 0300 2000 .....).
000000a0: 0000 1a2a 0000 b187 0200 1e00 0000 5a2a ...*.....Z*
000000b0: 0000 81a4 0200 0400 0000 6e61 6e00 0000 .....nan...
000000c0: 0000 b814 0000 b814 0000 b814 0000 b814 .....
000000d0: 0000 b814 0000 b814 0000 b814 0000 b814 .....
000000e0: 0000 b814 0000 b814 0000 b814 0000 b814 .....
000000f0: 0000 b814 0000 b814 0000 b814 0000 b814 .....
00000100: 0000 b814 0000 b814 0000 b814 0000 b814 .....
00000110: 0000 b814 0000 b814 0000 b814 0000 b814 .....
00000120: 0000 b814 0000 b814 0000 b814 0000 b814 .....
00000130: 0000 b814 0000 b814 0000 b814 0000 b814 .....
00000140: 0000 b814 0000 b814 0000 b814 0000 b814 .....
00000150: 0000 b814 0000 b814 0000 b814 0000 b814 .....
00000160: 0000 b814 0000 b814 0000 b814 0000 b814 .....
00000170: 0000 b814 0000 b814 0000 b814 0000 b814 .....
00000180: 0000 b814 0000 b814 0000 b814 0000 b814 .....
00000190: 0000 b814 0000 b814 0000 b814 0000 b814 .....
000001a0: 0000 b814 0000 b814 0000 b814 0000 b814 .....
000001b0: 0000 b814 0000 b814 0000 b814 0000 b814 .....
000001c0: 0000 b814 0000 b814 0000 b814 0000 b814 .....
```

PCA2\_54188401.tiff



**Eugene Cheipesh**

[echeipesh@azavea.com](mailto:echeipesh@azavea.com)

GitHub: @echeipesh

[www.azavea.com](http://www.azavea.com)





# Cloud Native Workflows


“Hey, let's not download the whole file every time.”



# Cloud Native

- Network is between storage and computation
- Storage is chunked and probably Key/Value
- Computation must scale
- Coordination is limited and risky





# C O G

CLOUD OPTIMIZED  
GEOTIFF



# Cloud Optimized GeoTiff

- Defines an internal structure of GeoTiffs that is optimized for streaming reading from blob storage like S3 or Google Cloud Storage.
- That is, it's best utilized for requests that use **HTTP GET Range Requests**







## Cloud optimized GeoTIFF ¶

### Definition

A cloud optimized GeoTIFF is a regular GeoTIFF file, aimed at being hosted on a HTTP file server, whose internal organization is friendly for consumption by clients issuing [HTTP GET range request](#) ("bytes: start\_offset-end\_offset" HTTP header).

It contains at its beginning the metadata of the full resolution imagery, followed by the optional presence of overview metadata, and finally the imagery itself. To make it friendly with streaming and progressive rendering, we recommand starting with the imagery of the smallest overview and finishing with the imagery of the full resolution level.

More formally, the structure of such a file is:

- TIFF / BigTIFF signature
- IFD ([⇒ Image File Directory](#)) of full resolution image
- Values of TIFF tags that don't fit inline in the IFD directory, such as [TileOffsets?](#), [TileByteCounts?](#) and GeoTIFF keys
- Optional: IFD (Image File Directory) of first overview (typically subsampled by a factor of 2), followed by the values of its tags that don't fit inline
- Optional: IFD (Image File Directory) of second overview (typically subsampled by a factor of 4), followed by the values of its tags that don't fit inline
- ...
- Optional: IFD (Image File Directory) of last overview (typically subsampled by a factor of  $2^N$ ), followed by the values of its tags that don't fit inline
- Optional: tile content of last overview level
- ...
- Optional: tile content of first overview level
- Tile content of full resolution image.

### Unspecified points

- size of tile: 256 or 512 pixels are typical however
- compression methods allowed. typically, no compression, DEFLATE or LZW can be used for lossless, or JPEG for lossy. (note that DEFLATE while more efficient than LZW can

### How to generate it with GDAL

Given an input dataset in tif with already generated internal or external overviews, a cloud optimized GeoTIFF can be generated with:

<https://trac.osgeo.org/gdal/wiki/CloudOptimizedGeoTIFF>





# Cloud Optimized GeoTIFF

An imagery format for cloud-native geospatial processing

[Introduction](#)[Why COG?](#)[How it works](#)[Get Started](#)[Implementations](#)

## About

A Cloud Optimized GeoTIFF (COG) is a regular GeoTIFF file, aimed at being hosted on a HTTP file server, with an internal organization that enables more efficient workflows on the cloud. It does this by leveraging the ability of clients issuing [HTTP GET range requests](#) to ask for just the parts of a file they need.

[Learn More](#)

<http://www.cogeo.org/>

# GeoTiff

“A GeoTIFF file is a TIFF 6.0 file, and inherits the file structure as described in the corresponding portion of the TIFF spec. All GeoTIFF specific information is encoded in several additional reserved TIFF tags, and contains no private Image File Directories (IFD's), binary structures or other private information invisible to standard TIFF readers.”

<http://geotiff.maptools.org/spec/geotiff2.4.html#2.4>



---

# **TIFF<sup>TM</sup>**

## **Revision 6.0**

**Final — June 3, 1992**

# TIFF

- Tagged Image File Format
- Binary image is separated out into “**segments**”
- Contains metadata in a sequence of “**tags**”
- Can contain metadata for **multiple images** in different “image file directories” (IDFs)





# GeoTiff Custom Tags

- GeoTiff adds custom tag GeoKeys
  - non-TIFF keys for CRS map space etc.
- This is how we know where the pixels are on the surface of the earth (the geospatial bit).

<http://geotiff.maptools.org/spec/geotiff6.html#6.1>



# GeoTiff Streaming Read

- **Step 1:** Fetch the header
- **Step 2:** Decide what segments you want
- **Step 3:** Read the segments



Internet Engineering Task Force (IETF)  
Request for Comments: 7233  
Obsoletes: [2616](#)  
Category: Standards Track  
ISSN: 2070-1721

R. Fielding, Ed.  
Adobe  
Y. Lafon, Ed.  
W3C  
J. Reschke, Ed.  
greenbytes  
June 2014

## **Hypertext Transfer Protocol (HTTP/1.1): Range Requests**

### **Abstract**

The Hypertext Transfer Protocol (HTTP) is a stateless application-level protocol for distributed, collaborative, hypertext information systems. This document defines range requests and the rules for constructing and combining responses to those requests.

### **Status of This Memo**

This is an Internet Standards Track document.

This document is a product of the Internet Engineering Task Force (IETF). It represents the consensus of the IETF community. It has received public review and has been approved for publication by the Internet Engineering Steering Group (IESG). Further information on Internet Standards is available in [Section 2 of RFC 5741](#).

Information about the current status of this document, any errata, and how to provide feedback on it may be obtained at <http://www.rfc-editor.org/info/rfc7233>.






# Request: BBOX Mean








**IFH**  
*Image File Header*

**IFD**  
*Image File Directory*



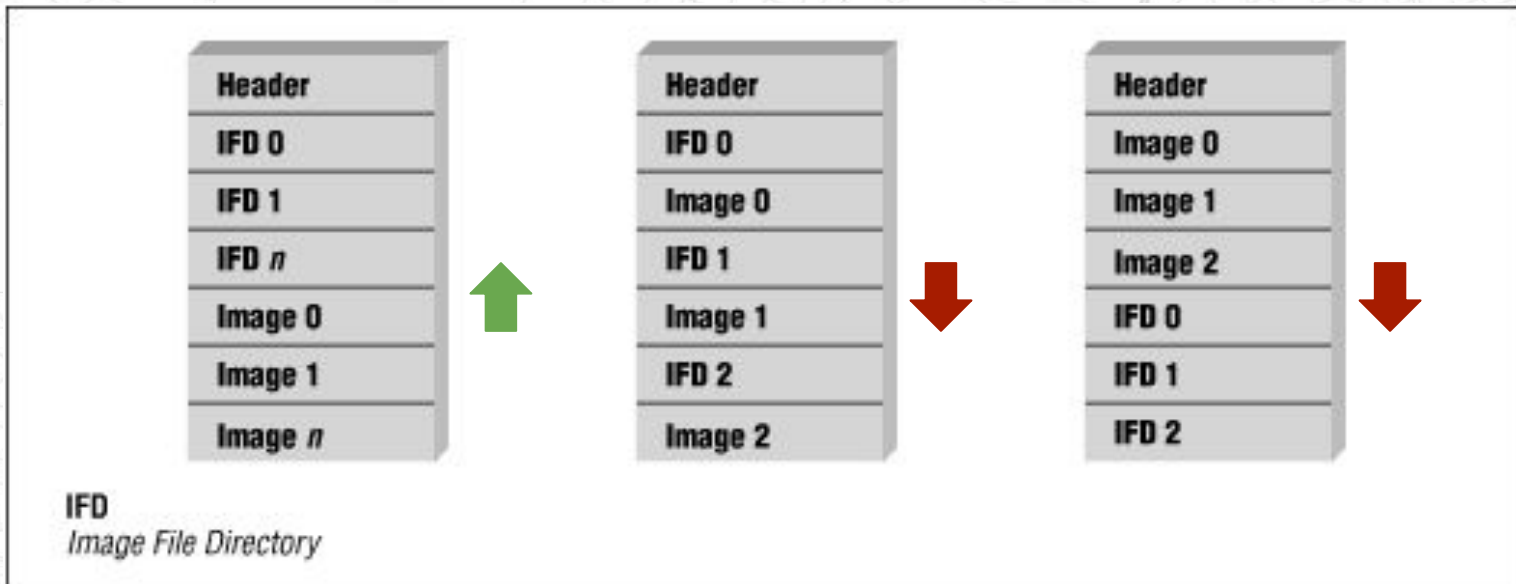
# GeoTIFF: All of the above



**IFD**  
*Image File Directory*



# COG: IFDs First - Rule #1




# TIFF Segments

- Chunks of an image are stored at different locations. This allows for parts of the image to be **decompressed** and loaded separately.
- There are two methods for segment layout: **Striped** and **Tiled**






# TIFF: Strip Segments



# COG: Strip Segments - Bad



# TIFF: Tile Segments





# COG: Use Tiled Segments - Rule #2







# Request: BBOX Mean at 60m



# Request: BBOX Mean at 60m




# GeoTIFF: Has Overviews





# COG: Provide Overviews - Rule #3





# COG Rules

- Put IFDs first
- Sort IFDs by resolution
- Use Tiled segments
- Include overviews (in file)



# COG Rules Continued ?

- Provenance Metadata
- Include Summary Statistics
- File naming convention
- Catalog suggestion



# COG “Profiles”: Slippy Map

- Restrict CRS as EPSG:3857
- Align GeoTiff segments with slippy tile grid
- Align GeoTiff files with tile boundaries
- Serve tile endpoints directly from COG



# COG World Coverage?

“You’re going to need more than one GeoTIFF.”





&lt;&gt; Code

! Issues 22

🔗 Pull requests 1

📁 Projects 0

📊 Insights

Branch: master ▾

stac-spec / json-spec /

Create new file

Upload files

Find file



mojdna Plural apostrophes

Latest commit f584425 5 da

..

examples	Plural apostrophes	5 da
json-schema	fix for time - added datetime property, but forgot to switch the requ...	2 mont
README.md	updated info on the EO profile	a mon
json-spec.md	changing start/end to datetime, as per #64	2 mont
package-lock.json	updated samples & examples to latest spec changes	2 mont
package.json	changed from underscore to dash	6 mont
sample-full.json	Update for latest EO profile	a mon
sample.json	updated samples & examples to latest spec changes	2 mont

📄 README.md

# STAC

## SpatioTemporal Asset Catalog



**s3://[bucket]/[Shard]\_[GeoHash]\_[ISO\_TIME]\_[ID].tiff**



# GeoTrellis & COGs








Scenes (0)

Add scenes


Import

Export





# COG: Tile Request



# COG Tiler

- Application specific catalog
- Reprojection on read
- Memache gives responsiveness
- Raster metadata pushed into files






# GeoTrellis Layers



# GeoTrellis: Avro Layers



→ Avro,  
SFC, → S3  
ZIP




# GeoTrellis Avro Layers


- Many small files, one per tile
  - PUT requests get expensive
- Data duplication hazzard
  - Probably not authoritative
- Limited access
  - Requires GeoTrellis code to read



# GeoTrellis: COG Layers



→ SFC → S3



EPSG 3857

Cells 31x31

Meta 4x4






# GeoTrellis Structured COG

- Base zoom tiles are in base GeoTiff layer
- Introduces “**DATE\_TIME**” tag
- Includes **VRT** for layer files
- GeoTiff **segments match** tile size and layout
- Additional zoom levels stored in **overviews**
- **Pyramid** up when overviews “run out”



# GeoTrellis: COG Pyramid










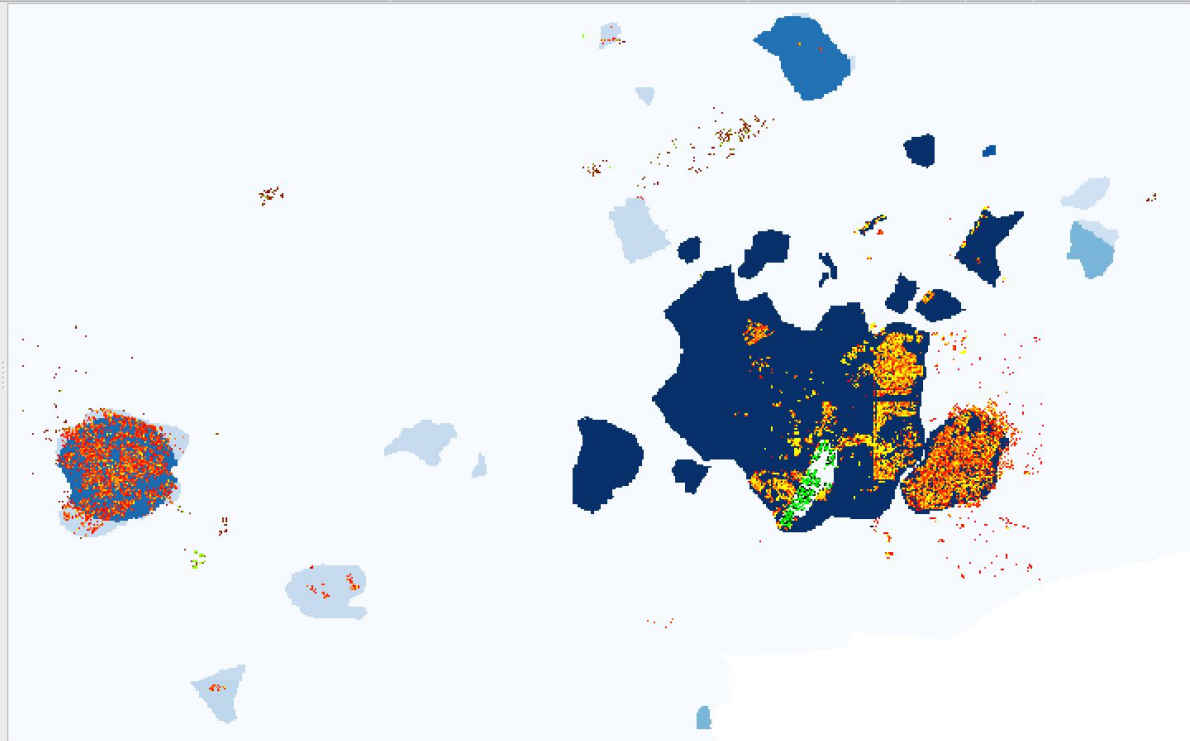
# HOTOSM - QGIS



## Layers



- ☒  **footprints**
- ☐  OSM Standard
- ☒  **label-base**
- ☐  lowpass-nan
- ☐  lowpass-zero
- ☐  Mapbox Satellite Streets
- ☒  **worldpop**



Q Type to locate (%%K)

2838543,-2823851



1:266,494



100%



0.0 °



☒ Render



EPSG:3857



# Unstructured COGs

Cloud Native thinking






# GeoTrellis Unstructured COG

- Lots of good raster sources are already COG
- They don't have a strict layout though
- Lets “inline” ingest process into the batch job
- Slower but reduces data duplication



# COG: Tile Request (Again)



# GeoTrellis Unstructured COG

- Requires mechanism to query for raster names
  - WFS, STAC, PostgreSQL, Scan
- COG provies “gradual” wins over GeoTiff
- Reproject on read to match workflow



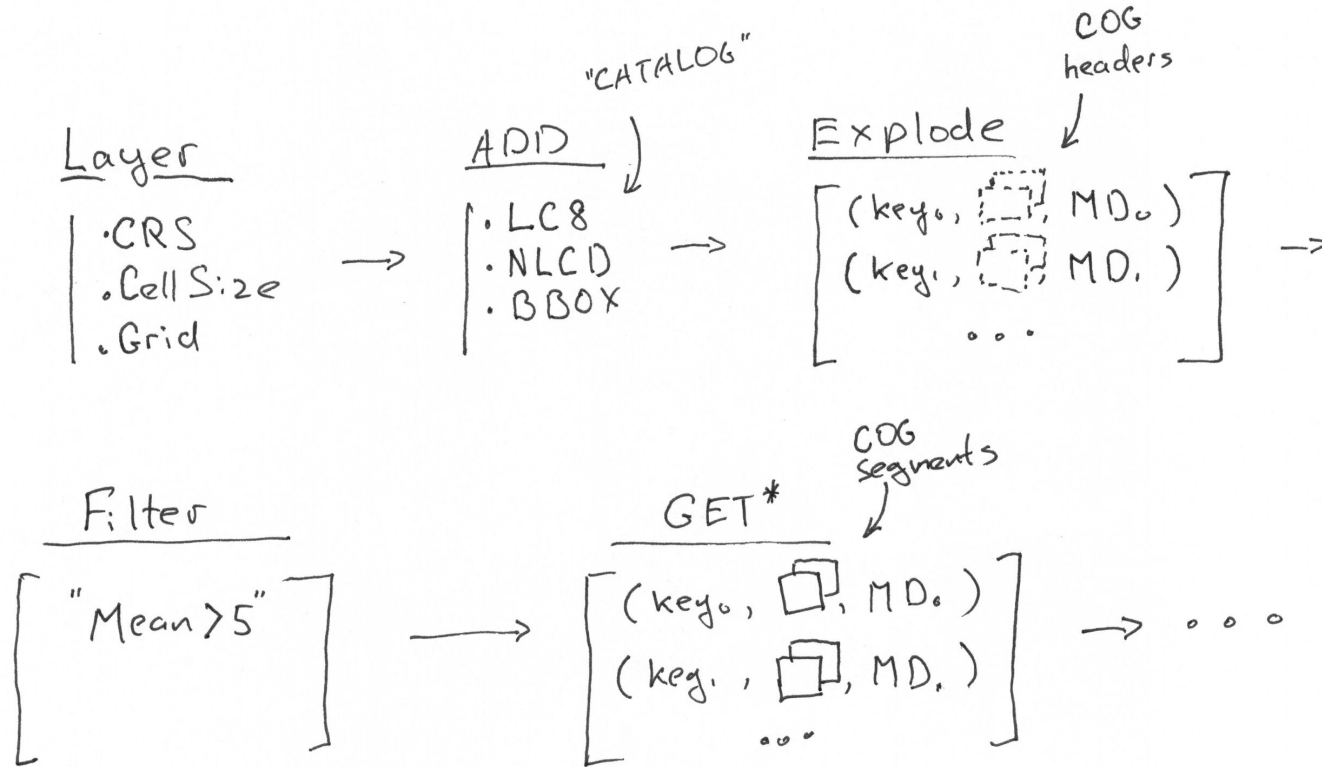
# Coregistration Workflow

- Working with **multiple** raster layers with differing:
  - CRS, Resolution, Extent
- Must pick the working/target layout
- We can “push-down” this process into IO





# Coregistration Workflow





# COGs for GeoTrellis

- Treating GeoTiff as a Key/Value tile store
  - With built in notion level of detail
- Turns an ingest workflow into query workflow
  - Use the metadata to plan the read
- Opens up the results
  - QGIS, GeoServer, GDAL



# COGs in General

- GeoJSON of raster formats
- Gradated spec
- Set of best practices
- Driven by implementations





# THANKS



**Q: Is this a good idea ?**



**Q: So is this a standard ?**



**Q: Why not use MRF ?**

